


The goal of this project is to develop a machine learning model for predicting customer churn in a telco company. Customer churn refers to the phenomenon where customers terminate their relationship with a company and switch to a competitor or discontinue using the service altogether. By analyzing various customer attributes and service usage patterns, we aim to create a classification model that can accurately predict whether a customer is likely to churn or not.

```
import pandas as pd
import numpy as np
import lightgbm as lgb
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.feature_selection import SelectKBest #for feature selecti
from sklearn.feature_selection import mutual_info_classif #for fear
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import StandardScaler,MinMaxScaler
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report,ConfusionMatrixDispla
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.tree import DecisionTreeClassifier
from sklearn.linear_model import LogisticRegression
from sklearn import tree
from sklearn.ensemble import RandomForestClassifier,StackingClassifier
from xgboost import XGBClassifier
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.ensemble import AdaBoostClassifier
from matplotlib import pyplot as plt
from sklearn.decomposition import PCA
from imblearn.over_sampling import SMOTE,RandomOverSampler
from sklearn.neural_network import MLPClassifier
```

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecur:
0	7590-VHVEG	Female	0	Yes	No	1	No	No phone service	DSL	
1	5575-GNVDE	Male	0	No	No	34	Yes	No	DSL	,
2	3668-QPYBK	Male	0	No	No	2	Yes	No	DSL	,
3	7795-CFOCW	Male	0	No	No	45	No	No phone service	DSL	,
4	9237-HQITU	Female	0	No	No	2	Yes	No	Fiber optic	
...	
7038	6840-RESVB	Male	0	Yes	Yes	24	Yes	Yes	DSL	,
7039	2234-XADUH	Female	0	Yes	Yes	72	Yes	Yes	Fiber optic	
7040	4801-JJAZL	Female	0	Yes	Yes	11	No	No phone service	DSL	,
7041	8361-LTMKD	Male	1	Yes	No	4	Yes	Yes	Fiber optic	
7042	3186-AJIEK	Male	0	No	No	66	Yes	No	Fiber optic	,

7043 rows × 21 columns



The dataset used for this project is a sample dataset of Telco Customer Churn, inspired by the original dataset "Telco customer churn (11.1.3+)" from IBM Business Analytics Community. It has been cleaned and aggregated for the purpose of telco customer churn analysis and prediction. The dataset is obtained from Kaggle.com.

```
df=pd.read_csv("/content/drive/MyDrive/data_luminar/WA_Fn-UseC_-Telco-Customer-Churn.csv")
df
```

The objective of this project is to build a classification model that can predict customer churn based on the given set of features. By analyzing the historical data and customer behavior patterns, the model will learn to identify the factors that contribute to churn and make accurate predictions for new customers. The ultimate aim is to assist the telco company in taking proactive measures to retain customers and improve customer satisfaction.

```
print("-----Visualizing the class Imbalance-----")
y=df["Churn"]
print(sns.countplot(data=df,y="Churn"))
df["Churn"].value_counts()
```

-----Visualizing the class Imbalance-----

Axes(0.125,0.11;0.775x0.77)

No 5174

Yes 1869

Name: Churn, dtype: int64



```
print(df.isna().sum())  
df.dtypes
```

```
customerID      0  
gender          0  
SeniorCitizen  0  
Partner         0  
Dependents      0  
tenure          0  
PhoneService    0  
MultipleLines   0  
InternetService 0  
OnlineSecurity  0  
OnlineBackup    0  
DeviceProtection 0  
TechSupport     0  
StreamingTV     0  
StreamingMovies 0  
Contract        0  
PaperlessBilling 0  
PaymentMethod   0  
MonthlyCharges  0  
TotalCharges    0  
Churn           0  
dtype: int64
```

customerID	object
gender	object
SeniorCitizen	int64
Partner	object
Dependents	object
tenure	int64
PhoneService	object
MultipleLines	object
InternetService	object
OnlineSecurity	object
OnlineBackup	object
DeviceProtection	object
TechSupport	object
StreamingTV	object
StreamingMovies	object
Contract	object
PaperlessBilling	object
PaymentMethod	object
MonthlyCharges	float64
TotalCharges	object
Churn	object
dtype:	object

Data Preprocessing: Perform data cleaning, handle missing values, and preprocess the categorical and numerical features as required.

```
df=df.drop(["customerID"],axis=1)
df['MultipleLines'] = df['MultipleLines'].replace('No phone service', 'No')
df
```

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	OnlineBackup
0	Female	0	Yes	No	1	No	No	DSL	No	
1	Male	0	No	No	34	Yes	No	DSL	Yes	
2	Male	0	No	No	2	Yes	No	DSL	Yes	
3	Male	0	No	No	45	No	No	DSL	Yes	
4	Female	0	No	No	2	Yes	No	Fiber optic	No	
...
7038	Male	0	Yes	Yes	24	Yes	Yes	DSL	Yes	
7039	Female	0	Yes	Yes	72	Yes	Yes	Fiber optic	No	
7040	Female	0	Yes	Yes	11	No	No	DSL	Yes	

```

le=LabelEncoder()
lst=["gender","Partner","OnlineSecurity","MonthlyCharges",
"Dependents",
"PhoneService",
"MultipleLines",
"InternetService",
"OnlineBackup",
"DeviceProtection",
"TechSupport",
"StreamingTV",
"StreamingMovies",

```

```
"Contract",  
"PaperlessBilling", "PaymentMethod", "TotalCharges", "Churn"]  
for i in lst:  
    df[i]=le.fit_transform(df[i])  
df.dtypes
```

```
gender          int64  
SeniorCitizen   int64  
Partner         int64  
Dependents      int64  
tenure          int64  
PhoneService    int64  
MultipleLines   int64  
InternetService int64  
OnlineSecurity  int64  
OnlineBackup    int64  
DeviceProtection int64  
TechSupport     int64  
StreamingTV     int64  
StreamingMovies int64  
Contract        int64  
PaperlessBilling int64  
PaymentMethod   int64  
MonthlyCharges  int64  
TotalCharges    int64  
Churn           int64  
dtype: object
```

```
df.corr()
```

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineS
gender	1.000000	-0.001874	-0.001808	0.010517	0.005106	-0.006488	-0.008414	-0.000863	-1
SeniorCitizen	-0.001874	1.000000	0.016479	-0.211185	0.016567	0.008576	0.142948	-0.032310	-1
Partner	-0.001808	0.016479	1.000000	0.452676	0.379697	0.017706	0.142057	0.000891	1
Dependents	0.010517	-0.211185	0.452676	1.000000	0.159712	-0.001762	-0.024526	0.044590	1
tenure	0.005106	0.016567	0.379697	0.159712	1.000000	0.008448	0.331941	-0.030359	1
PhoneService	-0.006488	0.008576	0.017706	-0.001762	0.008448	1.000000	0.279690	0.387436	-1
MultipleLines	-0.008414	0.142948	0.142057	-0.024526	0.331941	0.279690	1.000000	0.011124	1
InternetService	-0.000863	-0.032310	0.000891	0.044590	-0.030359	0.387436	0.011124	1.000000	-1
OnlineSecurity	-0.015017	-0.128221	0.150828	0.152166	0.325468	-0.015198	0.002306	-0.028416	1
OnlineBackup	-0.012057	-0.013632	0.153130	0.091015	0.370876	0.024105	0.119886	0.036138	1
DeviceProtection	0.000549	-0.021398	0.166330	0.080537	0.371105	0.003727	0.118577	0.044944	1
TechSupport	-0.006825	-0.151268	0.126733	0.133524	0.322942	-0.019158	0.005275	-0.026047	1
StreamingTV	-0.006421	0.030776	0.137341	0.046885	0.289373	0.055353	0.184681	0.107417	1
StreamingMovies	-0.008743	0.047266	0.129574	0.021321	0.296866	0.043870	0.186907	0.098350	1
Contract	0.000126	-0.142554	0.294806	0.243187	0.671607	0.002247	0.107114	0.099721	1
PaperlessBilling	-0.011754	0.156530	-0.014877	-0.111377	0.006152	0.016505	0.163530	-0.138625	-1
PaymentMethod	0.017352	0.028551	0.151708	0.040202	0.270126	0.004184	0.171026	0.096110	1

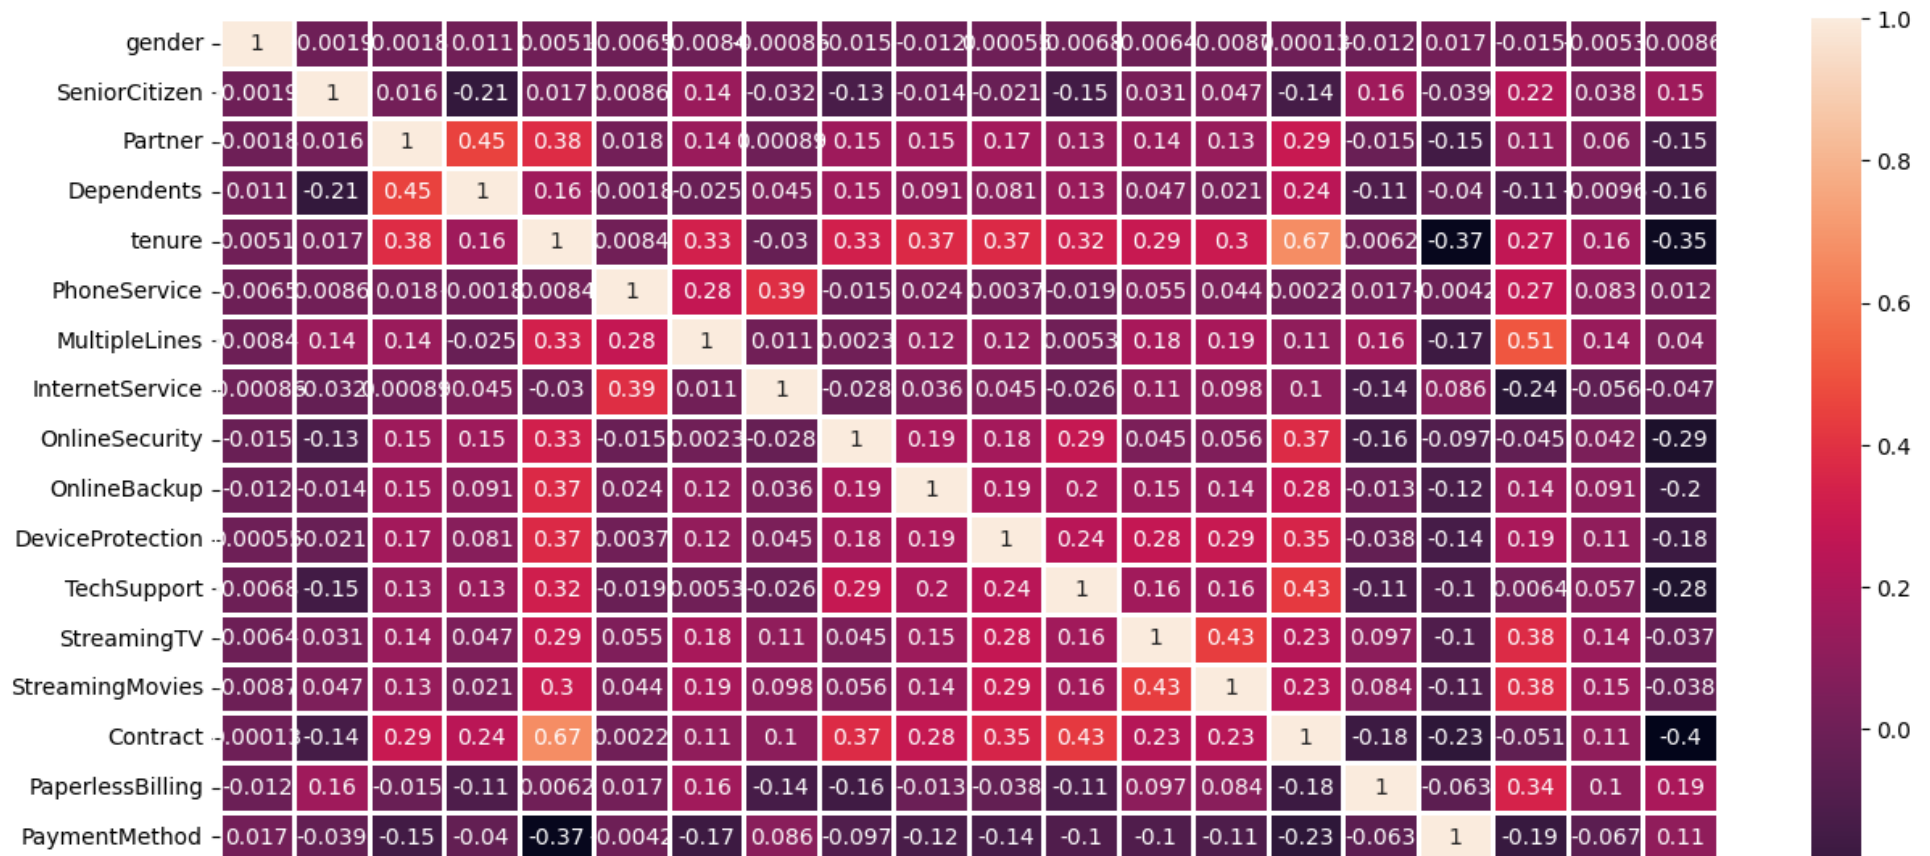
Exploratory Data Analysis: Analyze the dataset to gain insights into the distribution of features and their relationship with churn.

```
import warnings
print("-----finds potential relationships between variables and to understand the strength of these relationships")
plt.figure(figsize=[15,8])
```



```
warnings.filterwarnings("ignore")
plt.show(sns.heatmap(df.corr(),annot=True,linewidth=1))
```

-----finds potential relationships between variables and to understand the strength of these relationships.-----



```
df=df.drop(["gender"],axis=1)
```

```
df.dtypes
```

```
SeniorCitizen    int64
Partner          int64
Dependents       int64
tenure           int64
PhoneService     int64
```

```
MultipleLines      int64
InternetService    int64
OnlineSecurity      int64
OnlineBackup        int64
DeviceProtection    int64
TechSupport         int64
StreamingTV         int64
StreamingMovies     int64
Contract            int64
PaperlessBilling    int64
PaymentMethod       int64
MonthlyCharges      int64
TotalCharges        int64
Churn               int64
dtype: object
```

```
X=df.iloc[:, :-1]
y=df.iloc[:, -1]
```

```
#train_test_split
X_train,X_test,y_train,y_test=train_test_split(X,y,random_state=341,test_size=0.3)
print(X_train.shape[0]," rows is present in X_train")
```

```
4930 rows is present in X_train
```

```
# #feature selection
# bestfeatures=SelectKBest(score_func=mutual_info_classif,k=19)
# print("-----feature selection using the Information Gain method-----")
# scores = mutual_info_classif(X_train, y_train)
# # Create a list of feature names
# num_features = X_train.shape[1]
# feature_names = []
# for i in range(num_features):
```

```
# feature_names.append(f'Feature_{i}')

# df_scores = pd.DataFrame({'Feature':feature_names , 'Score': scores})
# df_scores = df_scores.sort_values('Score', ascending=False)
# print(df_scores)
# print("-----")
# num_columns = X.shape[1]
# column_names = [f'Feature_{i}' for i in range(num_columns)]
# column_name = X.columns[4]
# print(column_name)
```

▼ -----feature selection using the Information Gain method-----

Feature	Score
14 Feature_14	0.097724
4 Feature_4	0.071392
8 Feature_8	0.067328
11 Feature_11	0.057248
7 Feature_7	0.053660
16 Feature_16	0.048679
10 Feature_10	0.041474
17 Feature_17	0.039832
9 Feature_9	0.035735
12 Feature_12	0.032219
13 Feature_13	0.028845
15 Feature_15	0.023087
18 Feature_18	0.022439
3 Feature_3	0.015919
2 Feature_2	0.012065
6 Feature_6	0.011901
1 Feature_1	0.001936
0 Feature_0	0.001458
5 Feature_5	0.000346

Looking at the provided feature scores, ("gender") "Feature_0" has the lowest score of 0.002307. To potentially improve accuracy i consider to remove this particular column. even after removing features with low score accuracy is not improving so i dropped this process

```
#Scaling
sc=MinMaxScaler()
X_train=sc.fit_transform(X_train)
X_test=sc.transform(X_test)
```

Hyperparameter Tuning: Fine-tune the model's hyperparameters using techniques like grid search or random search to optimize its performance.

```
##HYPER PARAMETER TUNING
# from sklearn import feature_selection
# import warnings
# from sklearn.model_selection import GridSearchCV
# import warnings
# lgb_classifier=MLPClassifier()
# grid_vls={'hidden_layer_sizes': [(100,), (50, 50), (50, 25), (25, 25)],
#          'activation': ['relu', 'tanh'],
#          'solver': ['sgd', 'adam'],
#          'learning_rate': ['constant', 'adaptive']}
# grid=GridSearchCV(estimator=lgb_classifier, param_grid=grid_vls,scoring="accuracy", cv=5,refit=True,return_train_s
# grid.fit(X_train, y_train)
# grid.best_params_
```

Model Selection and Training: Experiment with various classification algorithms such as logistic regression, decision trees, random forests, or gradient boosting, and train the model using appropriate evaluation metrics.

normal method

```
#To execute machine learning models in a more efficient and organized manner,
#a recommended approach involves utilizing a for loop to iteratively run the models.

rf=RandomForestClassifier()
ad=AdaBoostClassifier()
nn = MLPClassifier(activation='tanh',hidden_layer_sizes= (100,),learning_rate="constant",solver="sgd")
```

```
lg=lgb.LGBMClassifier(learning_rate= 0.1,n_estimators=50,num_leaves=20)
lr=LogisticRegression()
gb=GradientBoostingClassifier()
xgb=XGBClassifier()
svc=SVC(C=0.1,gamma=0.01,kernel='linear',random_state=2)
dt=DecisionTreeClassifier(criterion="entropy",max_depth= 5,min_samples_leaf=4,min_samples_split=2)
lst1=[ad,dt,gb,rf,lr,svc,xgb,lg,nn]
ac=[]
model=[]
for i in lst1:
    print('*'*20,i,''*20)
    i.fit(X_train,y_train)
    y_pred=i.predict(X_test)
    print(classification_report(y_test,y_pred))

    print(i,ConfusionMatrixDisplay.from_predictions(y_test,y_pred))

    correct_predictions = (y_pred == y_test).sum()
    accuracy = correct_predictions / len(y_test)
    print("Accuracy:", accuracy)
    print("_"*200)
    ac.append(accuracy)
    model.append(i)
```

***** AdaBoostClassifier() *****

	precision	recall	f1-score	support
0	0.86	0.90	0.88	1576
1	0.66	0.55	0.60	537
accuracy			0.81	2113
macro avg	0.76	0.73	0.74	2113
weighted avg	0.81	0.81	0.81	2113

AdaBoostClassifier() <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay object at 0x7bdb86a606a0>
Accuracy: 0.8140085186938003

***** DecisionTreeClassifier(criterion='entropy', max_depth=5, min_samples_leaf=4) *****

	precision	recall	f1-score	support
0	0.84	0.89	0.86	1576
1	0.61	0.48	0.54	537
accuracy			0.79	2113
macro avg	0.72	0.69	0.70	2113
weighted avg	0.78	0.79	0.78	2113

DecisionTreeClassifier(criterion='entropy', max_depth=5, min_samples_leaf=4) <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay object at 0x7bdb8695d9c0>
Accuracy: 0.7893989588263133

***** GradientBoostingClassifier() *****

	precision	recall	f1-score	support
0	0.86	0.91	0.88	1576
1	0.67	0.55	0.61	537
accuracy			0.82	2113
macro avg	0.76	0.73	0.74	2113
weighted avg	0.81	0.82	0.81	2113

GradientBoostingClassifier() <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay object at 0x7bdb8695d9c0>
Accuracy: 0.8168480832938949

***** RandomForestClassifier() *****

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.84	0.90	0.87	1576
1	0.64	0.50	0.56	537
accuracy			0.80	2113
macro avg	0.74	0.70	0.72	2113
weighted avg	0.79	0.80	0.79	2113

RandomForestClassifier() <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay object at 0x7bdb869c4760>
Accuracy: 0.8021769995267393

***** LogisticRegression() *****

	precision	recall	f1-score	support
0	0.85	0.90	0.88	1576
1	0.65	0.55	0.60	537
accuracy			0.81	2113
macro avg	0.75	0.72	0.74	2113
weighted avg	0.80	0.81	0.80	2113

LogisticRegression() <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay object at 0x7bdb86824f70>
Accuracy: 0.8097491717936584

***** SVC(C=0.1, gamma=0.01, kernel='linear', random_state=2) *****

	precision	recall	f1-score	support
0	0.85	0.90	0.87	1576
1	0.65	0.54	0.59	537
accuracy			0.81	2113
macro avg	0.75	0.72	0.73	2113
weighted avg	0.80	0.81	0.80	2113

SVC(C=0.1, gamma=0.01, kernel='linear', random_state=2) <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay object at 0x7bdb86824f70>
Accuracy: 0.8078561287269286

***** XGBClassifier(base_score=None, booster=None, callbacks=None,
colsample_bylevel=None, colsample_bynode=None,
colsample_bytree=None, early_stopping_rounds=None,
enable_categorical=False, eval_metric=None, feature_types=None,

```

gamma=None, gpu_id=None, grow_policy=None, importance_type=None,
interaction_constraints=None, learning_rate=None, max_bin=None,
max_cat_threshold=None, max_cat_to_onehot=None,
max_delta_step=None, max_depth=None, max_leaves=None,
min_child_weight=None, missing=nan, monotone_constraints=None,
n_estimators=100, n_jobs=None, num_parallel_tree=None,
predictor=None, random_state=None, ...) *****
precision    recall  f1-score   support

```

```

0          0.85      0.88      0.87      1576
1          0.62      0.54      0.58       537

```

```

accuracy                0.80      2113
macro avg              0.73      0.71      0.72      2113
weighted avg           0.79      0.80      0.79      2113

```

```

XGBClassifier(base_score=None, booster=None, callbacks=None,
               colsample_bylevel=None, colsample_bynode=None,
               colsample_bytrees=None, early_stopping_rounds=None,
               enable_categorical=False, eval_metric=None, feature_types=None,
               gamma=None, gpu_id=None, grow_policy=None, importance_type=None,
               interaction_constraints=None, learning_rate=None, max_bin=None,
               max_cat_threshold=None, max_cat_to_onehot=None,
               max_delta_step=None, max_depth=None, max_leaves=None,
               min_child_weight=None, missing=nan, monotone_constraints=None,
               n_estimators=100, n_jobs=None, num_parallel_tree=None,
               predictor=None, random_state=None, ...) <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay object at
Accuracy: 0.7974443918599148

```

```

***** LGBMClassifier(n_estimators=50, num_leaves=20) *****
precision    recall  f1-score   support

```

```

0          0.85      0.90      0.88      1576
1          0.66      0.54      0.60       537

```

```

accuracy                0.81      2113
macro avg              0.76      0.72      0.74      2113
weighted avg           0.80      0.81      0.81      2113

```

```

LGBMClassifier(n_estimators=50, num_leaves=20) <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay object at 0x7bdb8
Accuracy: 0.812588736393753


```






```
***** MLPClassifier(activation='tanh', solver='sgd') *****
```



	precision	recall	f1-score	support
0	0.85	0.90	0.87	1576
1	0.64	0.52	0.57	537
accuracy			0.80	2113
macro avg	0.74	0.71	0.72	2113
weighted avg	0.79	0.80	0.80	2113



```
MLPClassifier(activation='tanh', solver='sgd') <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay object at 0x7bdb8  
Accuracy: 0.8045433033601515
```















DECISION TREE

```
#DECISION TREE
plt.figure(figsize=(5,8))
plt.title("Customer_Churn_Decision_Tree")
plt.show(tree.plot_tree(dt,filled=True,
feature_names=["gender",
```

```
"SeniorCitizen",  
"Partner",  
"Dependents",  
"tenure",  
"PhoneService",  
"MultipleLines",  
"InternetService",  
"OnlineSecurity",  
"OnlineBackup",  
"DeviceProtection",  
"TechSupport",  
"StreamingTV",  
"StreamingMovies",  
"Contract",  
"PaperlessBilling",  
"PaymentMethod",  
"MonthlyCharges",  
"TotalCharges",  
"Churn"],  
class_names=["no","yes"]  
)
```


$$\begin{pmatrix} 4930 \\ 2113 \end{pmatrix}, \begin{pmatrix} 3 \\ 3 \end{pmatrix}$$


```
rf=RandomForestClassifier(n_estimators=500,n_jobs=-1,max_depth=9,oob_score=True,random_state=25,max_samples=0.25,min
ad=AdaBoostClassifier()
nn = MLPClassifier(activation='tanh',hidden_layer_sizes= (100,),learning_rate="constant",solver="sgd")
lg=lgb.LGBMClassifier(learning_rate= 0.1,n_estimators=50,num_leaves=20)
lr=LogisticRegression(C=0.5,class_weight=None,dual=False,fit_intercept=True,max_iter=100,solver="saga",random_state=
gb=GradientBoostingClassifier()
xgb=XGBClassifier(learning_rate=0.01,max_depth=4,n_estimators=150)
svc=SVC(C=0.1,gamma=0.01,kernel='linear',random_state=2)
dt=DecisionTreeClassifier(criterion="entropy",max_depth= 5,min_samples_leaf=4,min_samples_split=2)
lst1=[dt,gb,rf,lr,svc,xgb,lg,nn]
acp=[]
modelp=[]
for p in lst1:
    print('*'*20,p,''*20)
    p.fit(X_trainm,y_train)
    y_pred=p.predict(X_testm)

    print(classification_report(y_test,y_pred))
    print(p,ConfusionMatrixDisplay.from_predictions(y_test,y_pred))

    correct_predictions = (y_pred == y_test).sum()
    accuracy = correct_predictions / len(y_test)
    print("_"*200)
    acp.append(accuracy)
    modelp.append(p)
```

```
***** DecisionTreeClassifier(criterion='entropy', max_depth=5, min_samples_leaf=4) *****
```

	precision	recall	f1-score	support
0	0.80	0.93	0.86	1576
1	0.61	0.32	0.42	537
accuracy			0.77	2113
macro avg	0.70	0.62	0.64	2113
weighted avg	0.75	0.77	0.75	2113

```
DecisionTreeClassifier(criterion='entropy', max_depth=5, min_samples_leaf=4) <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay object at 0x7bdb86705090>
```

```
***** GradientBoostingClassifier() *****
```

	precision	recall	f1-score	support
0	0.83	0.89	0.86	1576
1	0.59	0.47	0.53	537
accuracy			0.78	2113
macro avg	0.71	0.68	0.69	2113
weighted avg	0.77	0.78	0.77	2113

```
GradientBoostingClassifier() <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay object at 0x7bdb86705090>
```

```
***** RandomForestClassifier(max_depth=9, max_samples=0.25, min_samples_leaf=3,
                               n_estimators=500, n_jobs=-1, oob_score=True,
                               random_state=25) *****
```

	precision	recall	f1-score	support
0	0.82	0.91	0.86	1576
1	0.60	0.42	0.49	537
accuracy			0.78	2113
macro avg	0.71	0.66	0.68	2113
weighted avg	0.76	0.78	0.77	2113

```
RandomForestClassifier(max_depth=9, max_samples=0.25, min_samples_leaf=3,
                        n_estimators=500, n_jobs=-1, oob_score=True,
                        random_state=25) <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay object at 0x7bdb864d5b70>
```

```
***** LogisticRegression(C=0.5, random_state=15, solver='saga') *****
```

	precision	recall	f1-score	support
0	0.82	0.90	0.86	1576
1	0.59	0.42	0.49	537
accuracy			0.78	2113
macro avg	0.70	0.66	0.67	2113
weighted avg	0.76	0.78	0.76	2113

LogisticRegression(C=0.5, random_state=15, solver='saga') <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay object

***** SVC(C=0.1, gamma=0.01, kernel='linear', random_state=2) *****

	precision	recall	f1-score	support
0	0.82	0.89	0.86	1576
1	0.58	0.44	0.50	537
accuracy			0.78	2113
macro avg	0.70	0.67	0.68	2113
weighted avg	0.76	0.78	0.77	2113

SVC(C=0.1, gamma=0.01, kernel='linear', random_state=2) <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay object a

***** XGBClassifier(base_score=None, booster=None, callbacks=None,
 colsample_bylevel=None, colsample_bynode=None,
 colsample_bytrees=None, early_stopping_rounds=None,
 enable_categorical=False, eval_metric=None, feature_types=None,
 gamma=None, gpu_id=None, grow_policy=None, importance_type=None,
 interaction_constraints=None, learning_rate=0.01, max_bin=None,
 max_cat_threshold=None, max_cat_to_onehot=None,
 max_delta_step=None, max_depth=4, max_leaves=None,
 min_child_weight=None, missing=nan, monotone_constraints=None,
 n_estimators=150, n_jobs=None, num_parallel_tree=None,
 predictor=None, random_state=None, ...) *****

	precision	recall	f1-score	support
0	0.81	0.91	0.86	1576
1	0.60	0.39	0.47	537
accuracy			0.78	2113
macro avg	0.71	0.65	0.67	2113

weighted avg 0.76 0.78 0.76 2113

```
XGBClassifier(base_score=None, booster=None, callbacks=None,
              colsample_bylevel=None, colsample_bynode=None,
              colsample_bytree=None, early_stopping_rounds=None,
              enable_categorical=False, eval_metric=None, feature_types=None,
              gamma=None, gpu_id=None, grow_policy=None, importance_type=None,
              interaction_constraints=None, learning_rate=0.01, max_bin=None,
              max_cat_threshold=None, max_cat_to_onehot=None,
              max_delta_step=None, max_depth=4, max_leaves=None,
              min_child_weight=None, missing=nan, monotone_constraints=None,
              n_estimators=150, n_jobs=None, num_parallel_tree=None,
              predictor=None, random_state=None, ...) <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay object at
```

```
***** LGBMClassifier(n_estimators=50, num_leaves=20) *****
```

	precision	recall	f1-score	support
0	0.82	0.90	0.86	1576
1	0.60	0.43	0.50	537
accuracy			0.78	2113
macro avg	0.71	0.67	0.68	2113
weighted avg	0.77	0.78	0.77	2113


```
LGBMClassifier(n_estimators=50, num_leaves=20) <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay object at 0x7bdb8
```


```
***** MLPClassifier(activation='tanh', solver='sgd') *****
```



	precision	recall	f1-score	support
0	0.82	0.90	0.86	1576
1	0.58	0.41	0.48	537
accuracy			0.78	2113
macro avg	0.70	0.65	0.67	2113
weighted avg	0.76	0.78	0.76	2113



```
MLPClassifier(activation='tanh', solver='sgd') <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay object at 0x7bdb8
```














▼ SMOTE SAMPLING

```
sm=SMOTE(random_state=589)
X_train_o, y_train_o=sm.fit_resample(X_train,y_train)
#Finding best random state.
```

```
import random

random_numbers = random.sample(range(801,900),25)

best_random_state=None
best_accuracy=0.0
for value in random_numbers:
    X_train1,X_test1,y_train1,y_test1=train_test_split(X,y,test_size=0.3,random_state=value)
    sm=SMOTE()
    X_train, y_train=sm.fit_resample(X_train,y_train)
    for clf in lst1:
        clf.fit(X_train,y_train)
        accuracy=clf.score(X_test,y_test)
    if accuracy>0.80028395646000:    #best_accuracy    #589
        best_accuracy=accuracy
        best_random_state=value
        print("Best Random State:", best_random_state)
        print("Best Accuracy:", best_accuracy)
```

```
##feature selection
# from sklearn import feature_selection
# import warnings
# from sklearn.model_selection import GridSearchCV
# import warnings
# rf = XGBClassifier()
# grid_vls={'max_depth': [3, 5, 7],
#           'learning_rate': [0.1, 0.01, 0.001],
#           'n_estimators': [100, 200, 300],
#           'subsample': [0.8, 1.0],
#           'colsample_bytree': [0.8, 1.0]
#           }
# grid=GridSearchCV(estimator=rf, param_grid=grid_vls,scoring="accuracy", cv=5,refit=True,return_train_score=True)
```

```
# grid.fit(X_train_o, y_train_o)
# grid.best_params_

rf=RandomForestClassifier(max_depth=None,min_samples_leaf=1,min_samples_split=2,n_estimators=300)
ad=AdaBoostClassifier()

lg=lgb.LGBMClassifier(learning_rate= 0.1,n_estimators=50,num_leaves=20)
lr=LogisticRegression()
gb=GradientBoostingClassifier(learning_rate=0.05,max_depth= 5,n_estimators=150)
xgb=XGBClassifier(colsample_bytree=0.8,learning_rate= 0.1,max_depth=7,n_estimators=100,subsample=0.8)
svc=SVC()
dt=DecisionTreeClassifier(criterion="entropy",max_depth= 5,min_samples_leaf=4,min_samples_split=2)
nn = MLPClassifier(activation='tanh',hidden_layer_sizes= (100,),learning_rate="constant",solver="sgd")
lst1=[ad,dt,gb,rf,lr,svc,xgb,lg,nn]
ac_o=[]
model_o=[]
for j in lst1:
    print('*'*20,j,''*20)
    j.fit(X_train_o,y_train_o)
    y_pred=j.predict(X_test)

    print(classification_report(y_test,y_pred))
    print(j,ConfusionMatrixDisplay.from_predictions(y_test,y_pred))
    print("_"*200)
    correct_predictions = (y_pred == y_test).sum()
    accuracy = correct_predictions / len(y_test)

    ac_o.append(accuracy)
    # ac_o.append(j.accuracy_score(y_test,y_pred))
    # model_o.append(j)
```

```
***** AdaBoostClassifier() *****
```

	precision	recall	f1-score	support
0	0.91	0.79	0.84	1576
1	0.55	0.76	0.64	537
accuracy			0.78	2113
macro avg	0.73	0.78	0.74	2113
weighted avg	0.82	0.78	0.79	2113

```
AdaBoostClassifier() <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay object at 0x7bdb80a92e90>
```

```
***** DecisionTreeClassifier(criterion='entropy', max_depth=5, min_samples_leaf=4) *****
```

	precision	recall	f1-score	support
0	0.90	0.76	0.82	1576
1	0.51	0.76	0.61	537
accuracy			0.76	2113
macro avg	0.71	0.76	0.72	2113
weighted avg	0.80	0.76	0.77	2113

```
DecisionTreeClassifier(criterion='entropy', max_depth=5, min_samples_leaf=4) <sklearn.metrics._plot.confusion_matrix.ConfusionMat
```

```
***** GradientBoostingClassifier(learning_rate=0.05, max_depth=5, n_estimators=150) *****
```

	precision	recall	f1-score	support
0	0.88	0.83	0.86	1576
1	0.58	0.67	0.62	537
accuracy			0.79	2113
macro avg	0.73	0.75	0.74	2113
weighted avg	0.80	0.79	0.80	2113

```
GradientBoostingClassifier(learning_rate=0.05, max_depth=5, n_estimators=150) <sklearn.metrics._plot.confusion_matrix.ConfusionMa
```

```
***** RandomForestClassifier(n_estimators=300) *****
```

	precision	recall	f1-score	support
0	0.87	0.86	0.86	1576
1	0.60	0.62	0.61	537

accuracy			0.80	2113
macro avg	0.73	0.74	0.74	2113
weighted avg	0.80	0.80	0.80	2113

RandomForestClassifier(n_estimators=300) <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay object at 0x7bdb866b04f0>

***** LogisticRegression() *****

	precision	recall	f1-score	support
0	0.91	0.75	0.82	1576
1	0.52	0.79	0.62	537

accuracy			0.76	2113
macro avg	0.71	0.77	0.72	2113
weighted avg	0.81	0.76	0.77	2113

LogisticRegression() <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay object at 0x7bdb86288250>

***** SVC() *****

	precision	recall	f1-score	support
0	0.89	0.78	0.83	1576
1	0.53	0.73	0.61	537

accuracy			0.76	2113
macro avg	0.71	0.75	0.72	2113
weighted avg	0.80	0.76	0.78	2113

VC() <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay object at 0x7bdb866489d0>

***** XGBClassifier(base_score=None, booster=None, callbacks=None,
 colsample_bylevel=None, colsample_bynode=None,
 colsample_bytree=0.8, early_stopping_rounds=None,
 enable_categorical=False, eval_metric=None, feature_types=None,
 gamma=None, gpu_id=None, grow_policy=None, importance_type=None,
 interaction_constraints=None, learning_rate=0.1, max_bin=None,
 max_cat_threshold=None, max_cat_to_onehot=None,
 max_delta_step=None, max_depth=7, max_leaves=None,
 min_child_weight=None, missing=nan, monotone_constraints=None,
 n_estimators=100, n_jobs=None, num_parallel_tree=None,

```

predictor=None, random_state=None, ...) *****
precision    recall  f1-score   support

   0         0.88     0.84     0.86     1576
   1         0.59     0.65     0.62     537

 accuracy          0.80     2113
 macro avg         0.73     0.75     0.74     2113
 weighted avg         0.80     0.80     0.80     2113

GBClassifier(base_score=None, booster=None, callbacks=None,
             colsample_bylevel=None, colsample_bynode=None,
             colsample_bytree=0.8, early_stopping_rounds=None,
             enable_categorical=False, eval_metric=None, feature_types=None,
             gamma=None, gpu_id=None, grow_policy=None, importance_type=None,
             interaction_constraints=None, learning_rate=0.1, max_bin=None,
             max_cat_threshold=None, max_cat_to_onehot=None,
             max_delta_step=None, max_depth=7, max_leaves=None,
             min_child_weight=None, missing=nan, monotone_constraints=None,
             n_estimators=100, n_jobs=None, num_parallel_tree=None,
             predictor=None, random_state=None, ...) <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay object at 0x
***** LGBMClassifier(n_estimators=50, num_leaves=20) *****
precision    recall  f1-score   support

   0         0.88     0.84     0.86     1576
   1         0.58     0.67     0.62     537

 accuracy          0.79     2113
 macro avg         0.73     0.75     0.74     2113
 weighted avg         0.80     0.79     0.80     2113

GBMClassifier(n_estimators=50, num_leaves=20) <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay object at 0x7bdb866
***** MLPClassifier(activation='tanh', solver='sgd') *****
precision    recall  f1-score   support



   0         0.91     0.74     0.82     1576
   1         0.51     0.80     0.62     537



 accuracy          0.76     2113



```


```
macro avg      0.71      0.77      0.72      2113  
weighted avg   0.81      0.76      0.77      2113
```



```
LogisticRegression(activation='tanh', solver='sgd') <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay object at 0x7bdb864
```














```
#RandomOverSampler
oversampler = RandomOverSampler()
X_train_resampled, y_train_resampled = oversampler.fit_resample(X, y)

rf=RandomForestClassifier(max_depth=None,min_samples_leaf=1,min_samples_split=2,n_estimators=300)
ad=AdaBoostClassifier()

lg=lgb.LGBMClassifier(learning_rate= 0.1,n_estimators=50,num_leaves=20)
```

```
lr=LogisticRegression()
gb=GradientBoostingClassifier(learning_rate=0.05,max_depth= 5,n_estimators=150)
xgb=XGBClassifier(colsample_bytree=0.8,learning_rate= 0.1,max_depth=7,n_estimators=100,subsample=0.8)
svc=SVC()
dt=DecisionTreeClassifier(criterion="entropy",max_depth= 5,min_samples_leaf=4,min_samples_split=2)
nn = MLPClassifier(activation='tanh',hidden_layer_sizes= (100,),learning_rate="constant",solver="sgd")
lst1=[ad,dt,gb,rf,lr,svc,xgb,lg,nn]
ac_ro=[]
model_ro=[]
for j in lst1:
    print('*'*20,j,''*20)
    j.fit(X_train_resampled,y_train_resampled)
    y_pred=j.predict(X_test)

    print(classification_report(y_test,y_pred))
    #print(j,ConfusionMatrixDisplay.from_predictions(y_test,y_pred))
    print("_"*200)
    correct_predictions = (y_pred == y_test).sum()
    accuracy = correct_predictions / len(y_test)

    ac_ro.append(accuracy)
    model_ro.append(j)
```



over_sampling with PCA - it's accuracy is very low

UNDER_SAMPILNG - gives comparatively less accuracy than over sampling.

```
from imblearn.under_sampling import RandomUnderSampler
ra=RandomUnderSampler(random_state=1)
X_train_u, y_train_u=ra.fit_resample(X_train,y_train)
```

```
rf=RandomForestClassifier(n_estimators=500,n_jobs=-1,max_depth=9,oob_score=True,random_state=25,max_samples=0.25,min
ad=AdaBoostClassifier()

lg=lgb.LGBMClassifier()
lr=LogisticRegression()
gb=GradientBoostingClassifier()
xgb=XGBClassifier(learning_rate=0.01,max_depth=4,n_estimators=150)
svc=SVC(C=0.1,gamma=0.01,kernel='linear',random_state=2)
dt=DecisionTreeClassifier(criterion="entropy",max_depth= 5,min_samples_leaf=4,min_samples_split=2)
nn = MLPClassifier(activation='tanh',hidden_layer_sizes= (100,),learning_rate="constant",solver="sgd")
lst1=[ad,dt,gb,rf,lr,svc,xgb,lg,nn]
acus=[]
modelus=[]
for k in lst1:
    print('*'*20,k,''*20)
    k.fit(X_train_u,y_train_u)
    y_pred=k.predict(X_test)

    print(classification_report(y_test,y_pred))
```

```
print("_"*200)
correct_predictions = (y_pred == y_test).sum()
accuracy = correct_predictions / len(y_test)

acus.append(accuracy)

modelus.append(k)
```


accuracy	0.71	0.76	0.72	2113
macro avg	0.71	0.76	0.72	2113
weighted avg	0.81	0.76	0.77	2113

***** LGBMClassifier() *****

	precision	recall	f1-score	support
0	0.87	0.85	0.86	1576
1	0.59	0.64	0.61	537
accuracy			0.80	2113
macro avg	0.73	0.74	0.74	2113
weighted avg	0.80	0.80	0.80	2113

***** MLPClassifier(activation='tanh', solver='sgd') *****

	precision	recall	f1-score	support
0	0.92	0.74	0.82	1576
1	0.52	0.81	0.63	537
accuracy			0.76	2113
macro avg	0.72	0.77	0.72	2113
weighted avg	0.82	0.76	0.77	2113

pca under sampling

```
#pca under sampling
pca=PCA(n_components=3,random_state=3)
X_trainm=pca.fit_transform(X_train)
X_testm=pca.transform(X_test)
pca.explained_variance_ratio_

array([0.21219839, 0.13569808, 0.08801288])
```

```
rf=RandomForestClassifier(n_estimators=500,n_jobs=-1,max_depth=9,oob_score=True,random_state=25,max_samples=0.25,min
ad=AdaBoostClassifier()

lg=lgb.LGBMClassifier(learning_rate= 0.1,n_estimators=50,num_leaves=20)
lr=LogisticRegression()
#gb=GradientBoostingClassifier()
xgb=XGBClassifier(learning_rate=0.01,max_depth=4,n_estimators=150)
svc=SVC(C=0.1,gamma=0.01,kernel='linear',random_state=2)
dt=DecisionTreeClassifier(criterion="entropy",max_depth= 5,min_samples_leaf=4,min_samples_split=2)
nn = MLPClassifier(activation='tanh',hidden_layer_sizes= (100,),learning_rate="constant",solver="sgd")
lst1=[ad,dt,gb,rf,lr,svc,xgb,lg,nn]
acup=[]
modelup=[]
for up in lst1:
    print('*'*20,up,''*20)
    up.fit(X_trainm,y_train)
    y_pred=up.predict(X_testm)

    print(classification_report(y_test,y_pred))
    print("_"*200)
    correct_predictions = (y_pred == y_test).sum()
    accuracy = correct_predictions / len(y_test)

    acup.append(accuracy)

    modelup.append(up)
```



	precision	recall	f1-score	support
0	0.89	0.71	0.79	1576
1	0.47	0.75	0.58	537
accuracy			0.72	2113
macro avg	0.68	0.73	0.68	2113
weighted avg	0.79	0.72	0.74	2113

```

print(""*20,"stacking method",""*20)
lst2=[("knn",KNeighborsClassifier()),("nb",GaussianNB()),("ad",AdaBoostClassifier()),("dt",DecisionTreeClassifier())
stc=StackingClassifier(estimators=lst2,final_estimator=RandomForestClassifier())
stc.fit(X_train,y_train)
y_pred_stc=stc.predict(X_test)
print(classification_report(y_test,y_pred_stc))
ac_s=[]
model_s=[]
correct_predictions = (y_pred_stc == y_test).sum()
accuracy_stc = correct_predictions / len(y_test)
print("_"*200)

ac_s.append(accuracy_stc)
model_s.append(stc)
print(ac_s)

```

	precision	recall	f1-score	support
0	0.86	0.87	0.86	1576
1	0.60	0.57	0.58	537
accuracy			0.79	2113
macro avg	0.73	0.72	0.72	2113
weighted avg	0.79	0.79	0.79	2113

```
[0.7941315664931378]
```

```
from imblearn.combine import SMOTEENN
ou=SMOTEENN(random_state=1)
X_train_ou,y_train_ou=ou.fit_resample (X_train,y_train)
rf=RandomForestClassifier(n_estimators=500,n_jobs=-1,max_depth=9,oob_score=True,random_state=25,max_samples=0.25,min
ad=AdaBoostClassifier()

lg=lgb.LGBMClassifier(learning_rate= 0.1,n_estimators=50,num_leaves=20)
lr=LogisticRegression()
gb=GradientBoostingClassifier()
xgb=XGBClassifier(learning_rate=0.01,max_depth=4,n_estimators=150)
svc=SVC(C=0.1,gamma=0.01,kernel='linear',random_state=2)
dt=DecisionTreeClassifier(criterion="entropy",max_depth= 5,min_samples_leaf=4,min_samples_split=2)
nn = MLPClassifier(activation='tanh',hidden_layer_sizes= (100,),learning_rate="constant",solver="sgd")
lst1=[ad,dt,gb,rf,lr,svc,xgb,lg,nn]
acou=[]
modelou=[]
for up in lst1:
    print('*'*20,up,''*20)
    up.fit(X_train_ou,y_train_ou)
    y_pred=up.predict(X_test)

    print(classification_report(y_test,y_pred))
    print("_"*200)
    correct_predictions = (y_pred == y_test).sum()
    accuracy = correct_predictions / len(y_test)

    acou.append(accuracy)

    modelou.append(ou)
```


	precision	recall	f1-score	support
0	0.93	0.66	0.77	1576
1	0.46	0.85	0.60	537
accuracy			0.71	2113
macro avg	0.69	0.76	0.68	2113
weighted avg	0.81	0.71	0.73	2113

```

print("Highest accuracy in each method :-")
highAC_in_up=0.0
for m in acup:
    if m>highAC_in_up:
        highAC_in_up=m

highAC_in_ou=0.0
for m in acou:
    if m>highAC_in_ou:
        highAC_in_ou=m

highAC_in_normal=0.0
for m in ac:
    if m>highAC_in_normal:
        highAC_in_normal=m

highAC_in_us=0.0
for m in acus:
    if m>highAC_in_us:
        highAC_in_us=m

```



```

highAC_in_ro=0.0
for m in ac_ro:
    if m>highAC_in_ro:
        highAC_in_ro=m

highAC_in_o=0.0
for m in ac_o:
    if m>highAC_in_o:
        highAC_in_o=m

highAC_in_acp=0.0
for m in acp:
    if m>highAC_in_acp:
        highAC_in_acp=m

accuracy_stc
highAC_lst=pd.DataFrame({"undersampling_pca":highAC_in_up,"SMOTEENN":highAC_in_ou,"RandomOverSampling":highAC_in_ro,
print(highAC_lst)

```

Highest accuracy in each method :-

	undersampling_pca	SMOTEENN	RandomOverSampling	normal	undersampling \
0	0.729295	0.752958	0.769049	0.816848	0.795078

	smote	pca	stacking--
0	0.798391	0.78372	0.794132

A consolidated table was created to compare the accuracy of different algorithms and preprocessing techniques. The table was styled to highlight the highest accuracy value using red color and a gradient color map.

```
accuracy_SMOTEENN = pd.Series(acou, name='accuracy')
algorithm_SMOTEENN = pd.Series(modelou, name='algorithm')
normal = pd.concat([accuracy_SMOTEENN, algorithm_SMOTEENN], axis=1)

accuracy_normal = pd.Series(ac, name='accuracy')
algorithm_normal = pd.Series(model, name='algorithm')
normal = pd.concat([accuracy_normal, algorithm_normal], axis=1)

accuracy_RandomOverSampler = pd.Series(ac_ro, name='accuracy')
algorithm_RandomOverSampler = pd.Series(model_ro, name='algorithm')
RandomOverSampler = pd.concat([accuracy_RandomOverSampler, algorithm_RandomOverSampler], axis=1)

accuracy_pca = pd.Series(acp, name='accuracy')
algorithm_pca = pd.Series(modelp, name='algorithm')
pca = pd.concat([accuracy_pca, algorithm_pca], axis=1)

accuracy_under_sampling_pca = pd.Series(acup, name='accuracy')
algorithm_under_sampling_pca = pd.Series(modelup, name='algorithm')
under_sampling_pca = pd.concat([accuracy_under_sampling_pca, algorithm_under_sampling_pca], axis=1)

accuracy_over_smote = pd.Series(ac_o, name='accuracy')
algorithm_smote = pd.Series(model_o, name='algorithm')
smote = pd.concat([accuracy_over_smote, algorithm_smote], axis=1)

accuracy_under_sampling = pd.Series(acus, name='accuracy')
algorithm_under_sampling = pd.Series(modelus, name='algorithm')
under_sampling = pd.concat([accuracy_under_sampling, algorithm_under_sampling], axis=1)

aggregate_df=pd.concat([normal,pca,RandomOverSampler,under_sampling_pca,smote,under_sampling])
aggregate_df = aggregate_df.reset_index(drop=True) #because of difference in index number order highlight_max() di
```

```
aggregate_df_styled = aggregate_df.style.highlight_max(subset=['accuracy'], color="red")
aggregate_df_styled = aggregate_df_styled.background_gradient(subset=['accuracy'], cmap='RdYlGn', vmax=1.0, vmin=0.0)

print(" accuracy obtained through stacking",accuracy_stc)

aggregate_df_styled
```

accuracy obtained through stacking 0.7941315664931378

	accuracy	algorithm
0	0.814009	AdaBoostClassifier()
1	0.789399	DecisionTreeClassifier(criterion='entropy', max_depth=5, min_samples_leaf=4)
2	0.816848	GradientBoostingClassifier()
3	0.802177	RandomForestClassifier()
4	0.809749	LogisticRegression()
5	0.807856	SVC(C=0.1, gamma=0.01, kernel='linear', random_state=2)
6	0.797444	XGBClassifier(base_score=None, booster=None, callbacks=None, colsample_bylevel=None, colsample_bynode=None, colsample_bytree=None, early_stopping_rounds=None, enable_categorical=False, eval_metric=None, feature_types=None, gamma=None, gpu_id=None, grow_policy=None, importance_type=None, interaction_constraints=None, learning_rate=None, max_bin=None, max_cat_threshold=None, max_cat_to_onehot=None, max_delta_step=None, max_depth=None, max_leaves=None, min_child_weight=None, missing=nan, monotone_constraints=None, n_estimators=100, n_jobs=None, num_parallel_tree=None, predictor=None, random_state=None, ...)
7	0.812589	LGBMClassifier(n_estimators=50, num_leaves=20)
8	0.804543	MLPClassifier(activation='tanh', solver='sgd')
9	0.774255	DecisionTreeClassifier(criterion='entropy', max_depth=5, min_samples_leaf=4)
10	0.783720	GradientBoostingClassifier()
11	0.781827	RandomForestClassifier(max_depth=9, max_samples=0.25, min_samples_leaf=3, n_estimators=500, n_jobs=-1, oob_score=True, random_state=25)
12	0.777094	LogisticRegression(C=0.5, random_state=15, solver='saga')
13	0.778041	SVC(C=0.1, gamma=0.01, kernel='linear', random_state=2)
14	0.779934	XGBClassifier(base_score=None, booster=None, callbacks=None, colsample_bylevel=None, colsample_bynode=None, colsample_bytree=None, early_stopping_rounds=None, enable_categorical=False, eval_metric=None, feature_types=None, gamma=None, gpu_id=None, grow_policy=None, importance_type=None, interaction_constraints=None, learning_rate=0.01, max_bin=None, max_cat_threshold=None, max_cat_to_onehot=None, max_delta_step=None, max_depth=4, max_leaves=None, min_child_weight=None, missing=nan, monotone_constraints=None, n_estimators=150, n_jobs=None, num_parallel_tree=None, predictor=None, random_state=None, ...)

15	0.783247	LGBMClassifier(n_estimators=50, num_leaves=20)
16	0.775201	MLPClassifier(activation='tanh', solver='sgd')
17	0.745859	AdaBoostClassifier()
18	0.478467	DecisionTreeClassifier(criterion='entropy', max_depth=5, min_samples_leaf=4)
19	0.678183	GradientBoostingClassifier(learning_rate=0.05, max_depth=5, n_estimators=150)
20	0.604827	RandomForestClassifier(n_estimators=300)
21	0.769049	LogisticRegression()
22	0.745859	SVC()
23	0.761003	XGBClassifier(base_score=None, booster=None, callbacks=None, colsample_bylevel=None, colsample_bynode=None, colsample_bytree=0.8, early_stopping_rounds=None, enable_categorical=False, eval_metric=None, feature_types=None, gamma=None, gpu_id=None, grow_policy=None, importance_type=None, interaction_constraints=None, learning_rate=0.1, max_bin=None, max_cat_threshold=None, max_cat_to_onehot=None, max_delta_step=None, max_depth=7, max_leaves=None, min_child_weight=None, missing=nan, monotone_constraints=None, n_estimators=100, n_jobs=None, num_parallel_tree=None, predictor=None, random_state=None, ...)
24	0.718410	LGBMClassifier(n_estimators=50, num_leaves=20)
25	0.301467	MLPClassifier(activation='tanh', solver='sgd')
26	0.726455	AdaBoostClassifier()
27	0.703739	DecisionTreeClassifier(criterion='entropy', max_depth=5, min_samples_leaf=4)
28	0.722196	GradientBoostingClassifier()
29	0.729295	RandomForestClassifier(max_depth=9, max_samples=0.25, min_samples_leaf=3, n_estimators=500, n_jobs=-1, oob_score=True, random_state=25)
30	0.723142	LogisticRegression()
31	0.718410	SVC(C=0.1, gamma=0.01, kernel='linear', random_state=2)
32	0.722669	XGBClassifier(base_score=None, booster=None, callbacks=None, colsample_bylevel=None, colsample_bynode=None, colsample_bytree=None, early_stopping_rounds=None, enable_categorical=False, eval_metric=None, feature_types=None, gamma=None, gpu_id=None, grow_policy=None, importance_type=None, interaction_constraints=None, learning_rate=0.01, max_bin=None, max_cat_threshold=None, max_cat_to_onehot=None, max_delta_step=None, max_depth=4, max_leaves=None,

min_child_weight=None, missing=nan, monotone_constraints=None, n_estimators=150, n_jobs=None, num_parallel_tree=None, predictor=None, random_state=None, ...)

33 0.726455

LGBMClassifier(n_estimators=50, num_leaves=20)

34 0.720303

MLPClassifier(activation='tanh', solver='sgd')

35 0.781827

nan

36 0.756744

nan

37 0.792239

nan

38 0.798391

nan

39 0.758164

nan

40 0.764316

nan

41 0.795551

nan

42 0.766300

nan

